

# Understanding Kafka Connect

Intern Names : Parth.V.Sharma & Nikhil Mehta  
Mentor Name : Manas Joshi  
Manager Name : Aniket Khandelwal  
Team : Platform

# Project Overview :

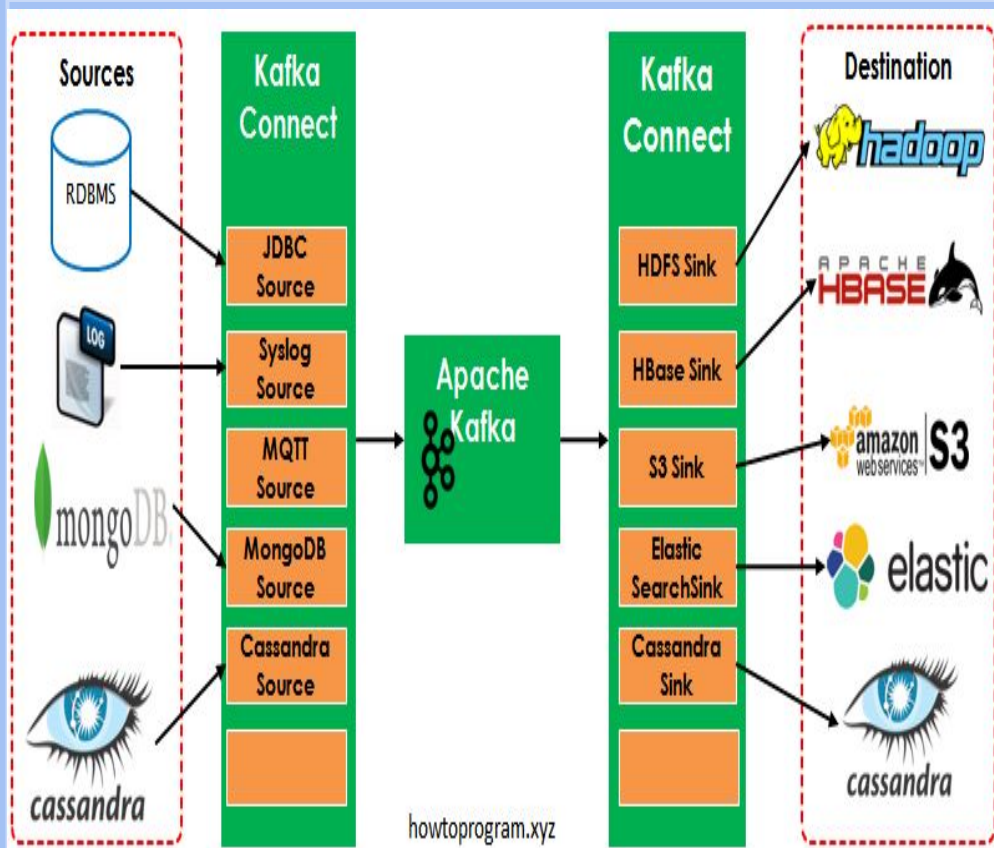
## Area of Study:

- What is Kafka Connect
- Provide Customization Features and explore flexibility
- Advantages

## Major Achievements:

- ➔ Data pipeline to move Data from **MongoDb** to **Redis** and/or **ElasticSearch**
- ➔ Created mechanisms for Error handling
- ➔ Provided Capability to modify data flowing in the pipeline

# What is Kafka Connect?



- ❏ A tool for scalably and reliably streaming data between Apache Kafka and other data systems
- ❏ Made up of a set of connectors which act as the medium of interaction between kafka and external systems


# Types of Kafka Connectors:

## Source Connector

1. Imports data from any external system into an Apache Kafka topic

## Sink Connector

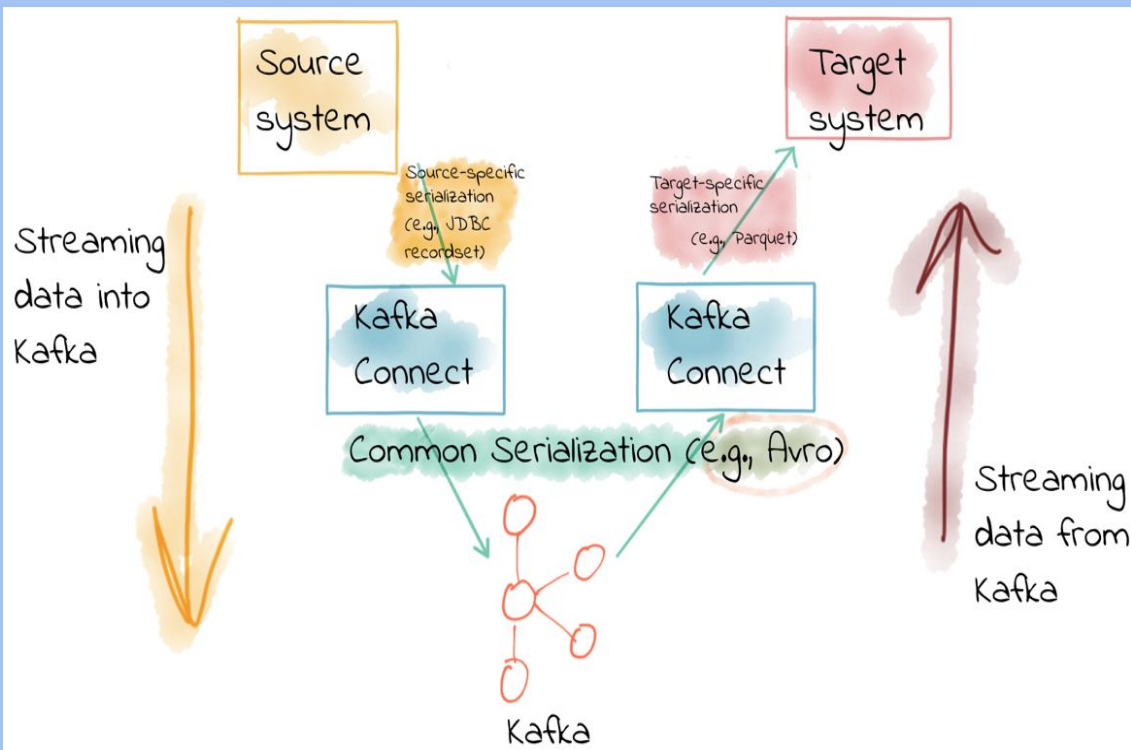
1. Sink connector allows you to export data from Apache Kafka topics to any other system



# Important Kafka Connect Features & Concepts

# Convertors

## Block Diagram




## Description:

- Handle serialization and deserialization of data
- Data is in the form of bytes when stored in Kafka
- Standardized serialization format
  - Json
  - Avro
  - String

# Single Message Transforms

## Block Diagram



## Description:

- Provides a way to Manipulate data when in flows in the pipeline
- Transform inbound messages after a source connector has produced them
- SMTs transform outbound messages before they are sent to a sink connector

# Single Message Transforms

## Some Useful SMTs

- **ValueToKey**
  - Replace the record key with a new key formed from a subset of fields in the record value.
- **ReplaceField**
  - Filter or rename fields.
- **Add - Drop Fields**
- **AlterSchema**
  - Change the schema of the key or value of the record
- **Custom SMT**



# Custom Single Message Transforms

Manipulate all parts of the Record: the Key, the Value, the Key and Value schemas, destination topic, destination partition, and timestamp

```
public class CustomSMT implements Transformation
```

## Configure method

Determines the input parameters to be defined in the connector configuration and properties regarding them

- ☐ Default Value
- ☐ Importance
- ☐ Validators

## Apply method

Takes a kafka record as input, performs certain operations and returns a modified kafka record to be inserted

```
@Override  
public R apply(R record) {}
```

# Custom SMT Examples

## Customized Dead Letter Queue

Store various exceptions that occur in the pipeline without stopping the connector

- Include more exceptions than the already existing ones
- Different DLQs for different exceptions

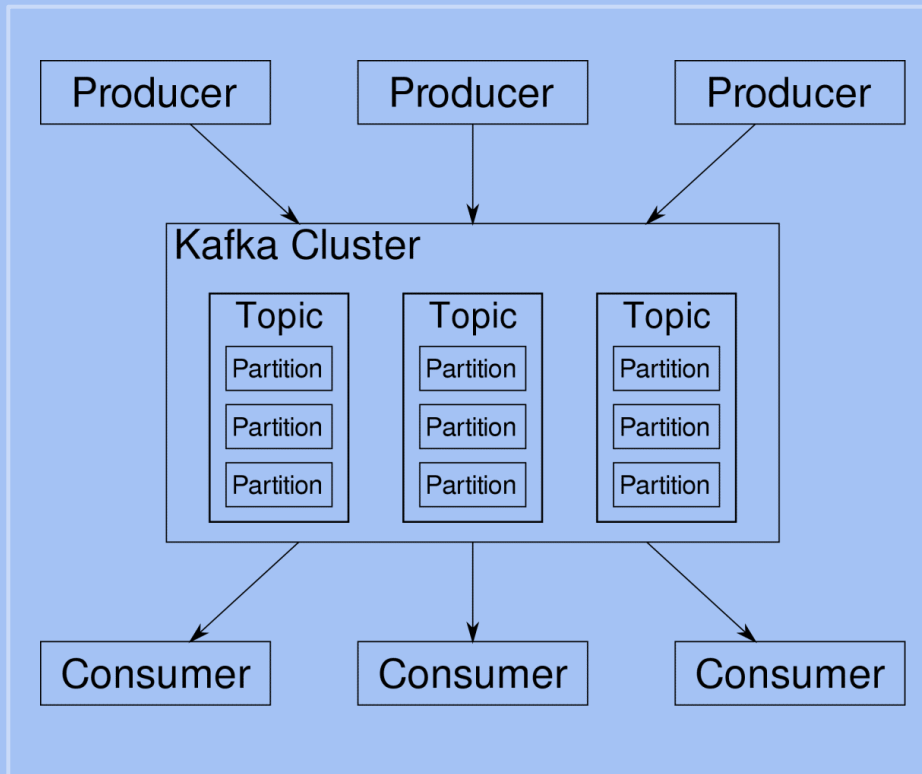
## Tracking records

Keep a count of :

- I. messages flowing out from source connector
- II. Messages flowing into sink connector

Ensure they are equal

# Kafka Partitioner & Partition Key



## Usage and Explanation

- ❑ Kafka partitioner is responsible for deciding partition number for each message.
  
- ❑ The default partitioner follows these rules.
  1. If a partition number is explicitly defined straightaway use it
  2. Else if it uses the partition key to choose a partition based on a hash value of this key
    - a. eg)  $\text{hash}(\text{key}) \% \text{num\_partition}$ .
  3. If no partition number or key is present, pick a partition in a round-robin fashion.

# Creating Custom Partitioner

## What?

- ❑ Create our own partitioning logic for the messages flowing from the source connector into a kafka topic

## Why?

- Default partitioner does not cater to some scenarios
- Data from same producer to go to the same partition
  - ◆ Using Composite key to partition data
  - ◆ The way Hashing works
  - ◆ Number of partitions increase
  - ◆ Reserving some partitions

## How?

- Kafka provides an interface called Partitioner
- custom partitioner class must implement three methods from the interface
  - ◆ `public void configure(Map<String, ?> configs)`
  - ◆ `public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[] valueBytes, Cluster cluster)`
  - ◆ `public void close()`

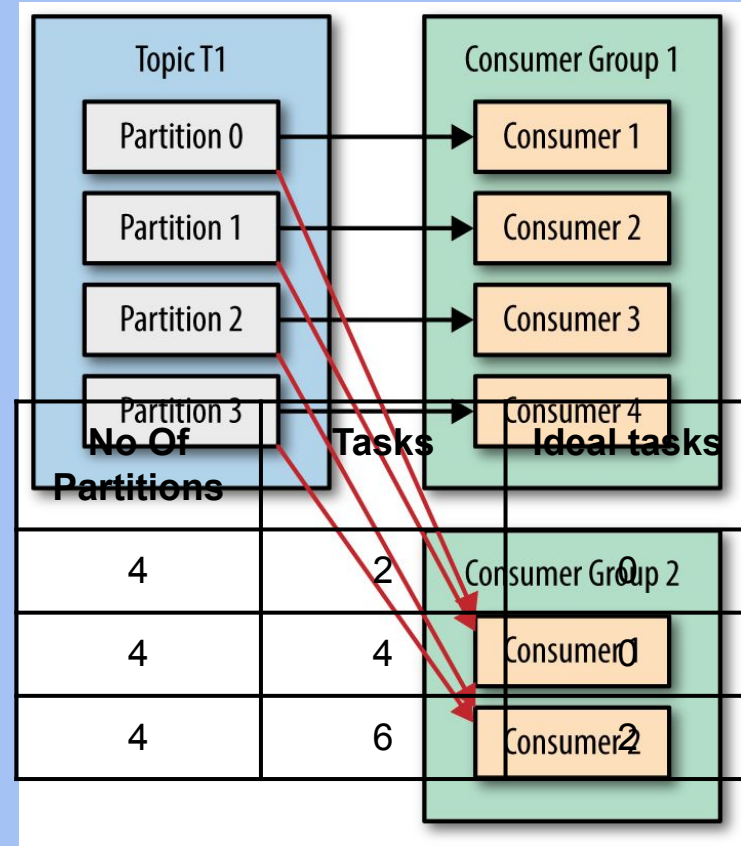
**Any Questions ?**

**Thank You**

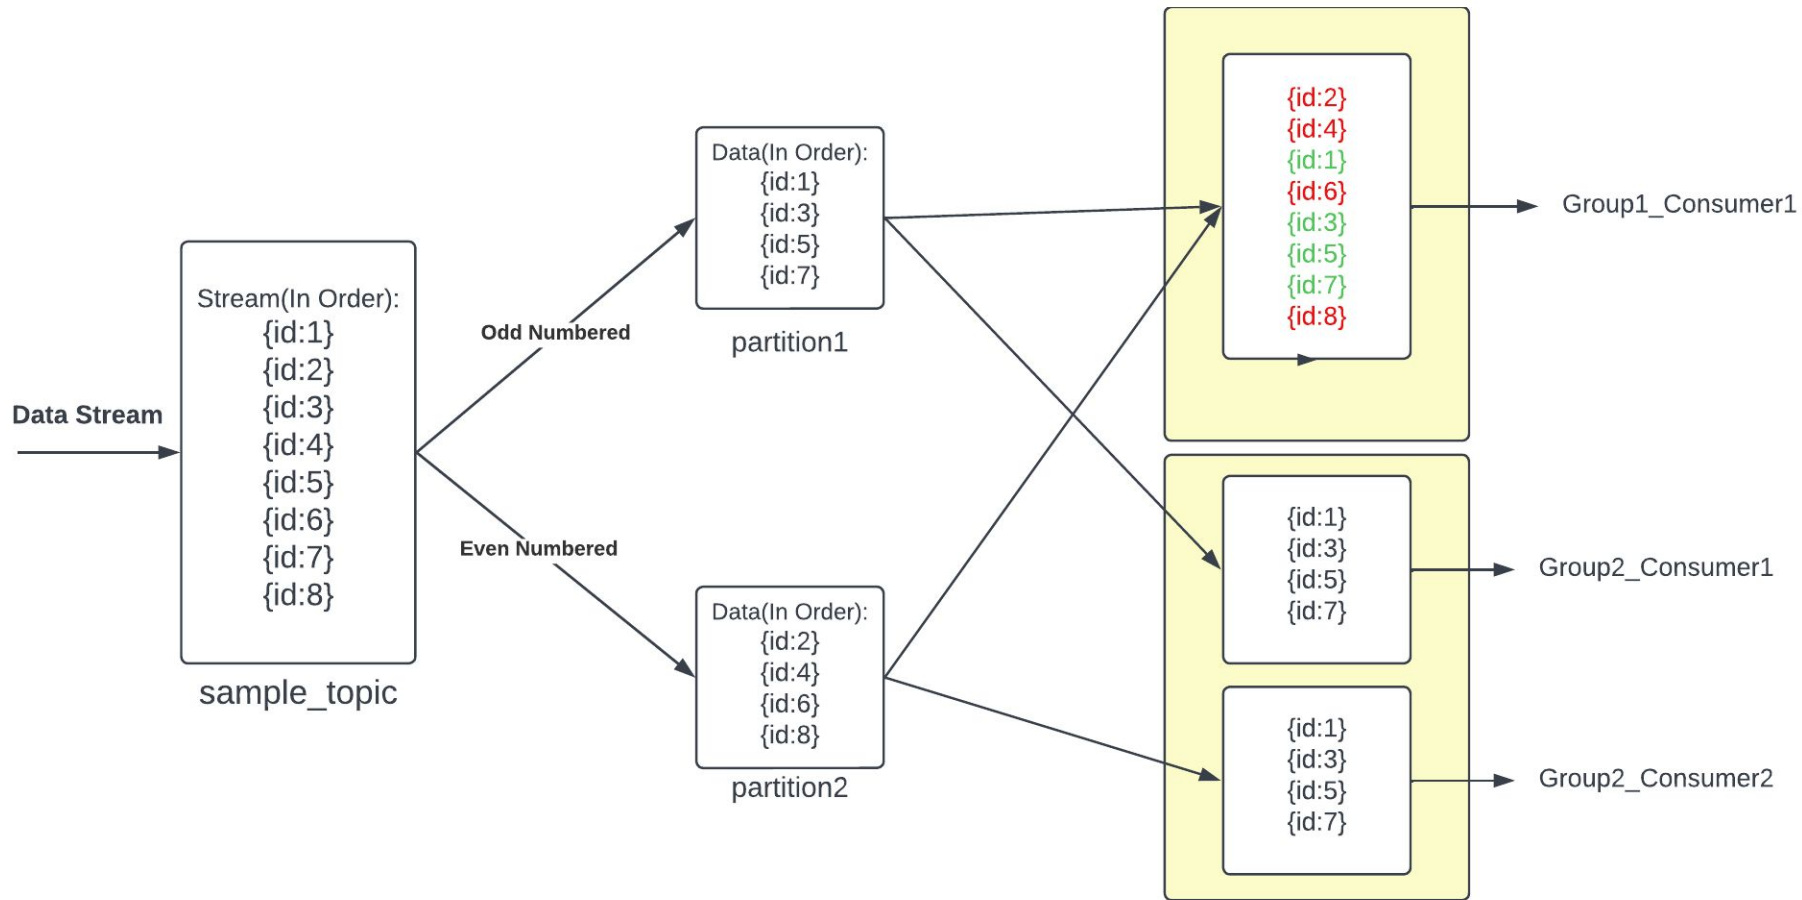
# Message Ordering

- Messages coming to a particular topic can be divided into partitions through partition key.
- Only one consumer from a group can subscribe to a partition.
- Consumers in different groups can subscribe to same partition.
- A consumer can subscribe to multiple partitions.

**Unit of Parallelism : Partition**



Source: [Oreilly](#)



# Error Handling



# Error Handling By Kafka Connect

Stage	Handled?
Start	NO
poll(source connector)	NO
convert	YES
transform	YES
put(sink connector)	NO

# Error Handling By Kafka Connect

Fail Fast	Ignore	Dead Letter Queue(DLQ)
<code>tolerance = none (default)</code>	<code>tolerance = all</code> <code>log.enable = true</code> <code>log.include.messages = true</code>	<code>tolerance = all</code> <code>deadletterqueue.topic.name = &lt;name&gt;</code> <code>deadletterqueue.context.headers.enabled = true</code>

DLQ topic can be subscribed by other consumers or if the error can be rectified, can be sent to same producer.

## No DLQ In Put Stage?

**Sink Database may be down.**

- We can set number of retries to a high number with high backoff interval.
- Use custom logic such as retries for particular error and DLQ for others.
- ES Sink Connector uses ‘\_bulk’ api and error could be due to one erroneous record.

# Error Handling in Put Stage

We can analyse connectors source code and apply custom error handling strategies.

Redis Connector	ES Connector
<pre>if (null == record.key()) {     throw new DataException("The key for the         record cannot be null. " +         <i>formatLocation</i>(record) };</pre>	<pre>if (shouldSkipRecord(record)) {     logTrace("Ignoring {} with null value.",         record);     offsetState.markProcessed();     reportBadRecord(record, new         ConnectException("Cannot write null valued             record."));     continue; }</pre>

# Advantages of Kafka connect

The benefits of Kafka Connect include:

- **Data Centric Pipeline** – Uses meaningful data abstractions
- **Flexibility and Scalability** – Standalone or Distributed
- **Reusability and Extensibility** – Connect leverages existing connectors or extends them to tailor to your needs and provides lower time to production.
- **Schema Registry** – Confluent provides a layer to manage schemas and schema evolution.

**Any Questions ?**

**Thank You**

# Learnings :

1. Java as a language for development .
2. Java Concurrency .
3. Distributed Systems .
4. Kafka Apache .
5. Implementation of Kafka building basic group chat api .
6. Unstructured Databases Redis , Scylla .
7. Understanding and reading of source codes .
8. Debugging and Fixing Bugs .
9. Jedis Java Api .
10. Methods of benchmarking .
11. Writing clean and extendible codes .
12. Using various tools on Git.
13. Getting familiar with jar files , Config files etc.
14. Using various helpful tools , docker , homebrew , throttle etc .